



## Chapter 11

# Database Performance Tuning and Query Optimization

# Learning Objectives

- In this chapter, you will learn:
  - Basic database performance-tuning concepts
  - How a DBMS processes SQL queries
  - About the importance of indexes in query processing
  - About the types of decisions the query optimizer has to make
  - Some common practices used to write efficient SQL code
  - How to formulate queries and tune the DBMS for optimal performance

# Database Performance-Tuning Concepts

- Goal of database performance is to execute queries as fast as possible
- **Database performance tuning:** Set of activities and procedures that reduce response time of database system
- Fine-tuning the performance of a system requires that all factors must operate at optimum level with minimal bottlenecks

# Table 11.1 - General Guidelines for Better System Performance

TABLE 11.1

**GENERAL GUIDELINES FOR BETTER SYSTEM PERFORMANCE**

	SYSTEM RESOURCES	CLIENT	SERVER
<b>Hardware</b>	CPU	The fastest possible Dual-core CPU or higher	The fastest possible Multiple processors (quad-core technology) Cluster of networked computers
	RAM	The maximum possible to avoid OS memory to disk swapping	The maximum possible to avoid OS memory to disk swapping
	Hard disk	Fast SATA/EIDE hard disk with sufficient free hard disk space Solid State Drives (SSD) for faster speed	Multiple high-speed, high-capacity disks Fast disk interface (SAS / SCSI / Firewire / Fibre Channel) RAID configuration optimized for throughput Solid State Drives (SSD) for faster speed Separate disks for OS, DBMS, and data spaces
	Network	High-speed connection	High-speed connection
<b>Software</b>	Operating System (OS)	64-bit OS for larger address spaces Fine-tuned for best client application performance	64-bit OS for larger address spaces Fine-tuned for best server application performance
	Network	Fine-tuned for best throughput	Fine-tuned for best throughput
	Application	Optimize SQL in client application	Optimize DBMS server for best performance

# Performance Tuning: Client and Server

- Client side
  - **SQL performance tuning:** Generates SQL query that returns correct answer in least amount of time
    - Using minimum amount of resources at server
- Server side
  - **DBMS performance tuning:** DBMS environment configured to respond to clients' requests as fast as possible
    - Optimum use of existing resources

# DBMS Architecture

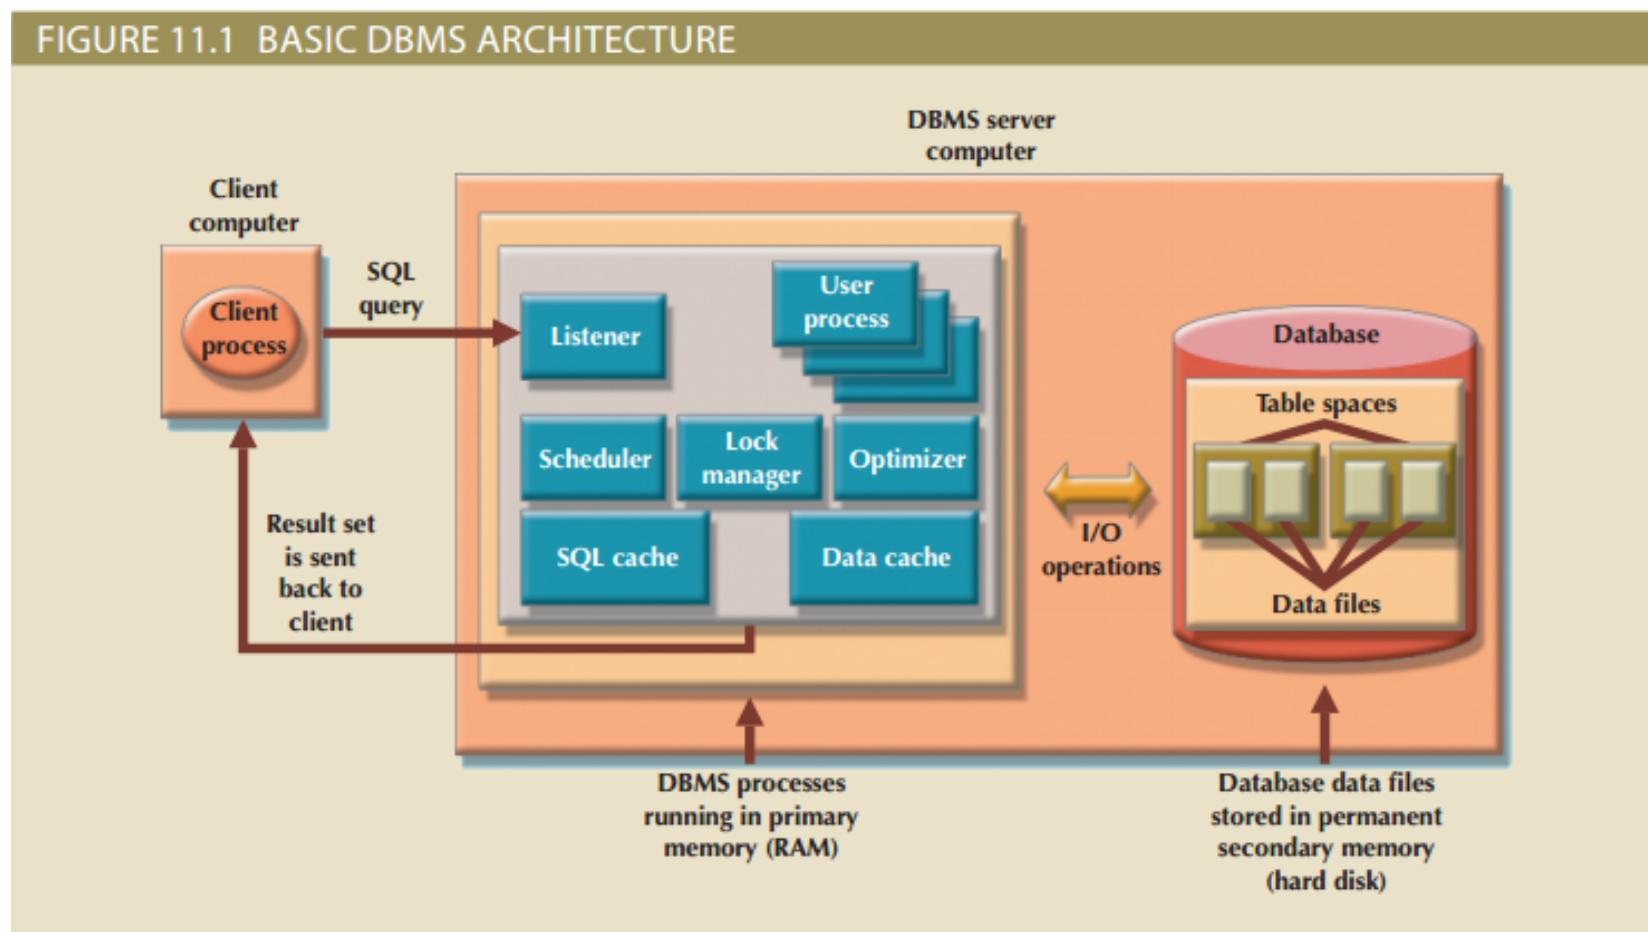
- All data in a database are stored in **data files**
  - Data files automatically expand in predefined increments known as **extends**
- Data files are grouped in file groups or table spaces
  - **Table space or file group:** Logical grouping of several data files that store data with similar characteristics
- **Data cache or buffer cache:** Shared, reserved memory area
  - Stores most recently accessed data blocks in RAM

# DBMS Architecture

- **SQL cache or procedure cache:** Stores most recently executed SQL statements or PL/SQL procedures
- DBMS retrieves data from permanent storage and places them in RAM
- **Input/output request:** Low-level data access operation that reads or writes data to and from computer devices
- Data cache is faster than working with data files
- Majority of performance-tuning activities focus on minimizing I/O operations

# Figure 11.1 - Basic DBMS Architecture

FIGURE 11.1 BASIC DBMS ARCHITECTURE



# Database Query Optimization Modes

- Algorithms proposed for query optimization are based on:
  - Selection of the optimum order to achieve the fastest execution time
  - Selection of sites to be accessed to minimize communication costs
- Evaluated on the basis of:
  - Operation mode
  - Timing of its optimization

# Classification of Operation Modes

- **Automatic query optimization:** DBMS finds the most cost-effective access path without user intervention
- **Manual query optimization:** Requires that the optimization be selected and scheduled by the end user or programmer

# Classification Based on Timing of Optimization

- **Static query optimization:** best optimization strategy is selected when the query is compiled by the DBMS
  - Takes place at compilation time
- **Dynamic query optimization:** Access strategy is dynamically determined by the DBMS at run time, using the most up-to-date information about the database
  - Takes place at execution time

# Classification Based on Type of Information Used to Optimize the Query

- **Statistically based query optimization algorithm:** Statistics are used by the DBMS to determine the best access strategy
- Statistical information is generated by DBMS through:
  - **Dynamic statistical generation mode**
  - **Manual statistical generation mode**
- **Rule-based query optimization algorithm:** based on a set of user-defined rules to determine the best query access strategy

# Table 11.2 - Sample Database Statistics Measurements

TABLE 11.2

## SAMPLE DATABASE STATISTICS MEASUREMENTS

DATABASE OBJECT	SAMPLE MEASUREMENTS
Tables	Number of rows, number of disk blocks used, row length, number of columns in each row, number of distinct values in each column, maximum value in each column, minimum value in each column, and columns that have indexes
Indexes	Number and name of columns in the index key, number of key values in the index, number of distinct key values in the index key, histogram of key values in an index, and number of disk pages used by the index
Environment Resources	Logical and physical disk block size, location and size of data files, and number of extends per data file

# Query Processing

## Parsing

- DBMS parses the SQL query and chooses the most efficient access/execution plan

## Execution

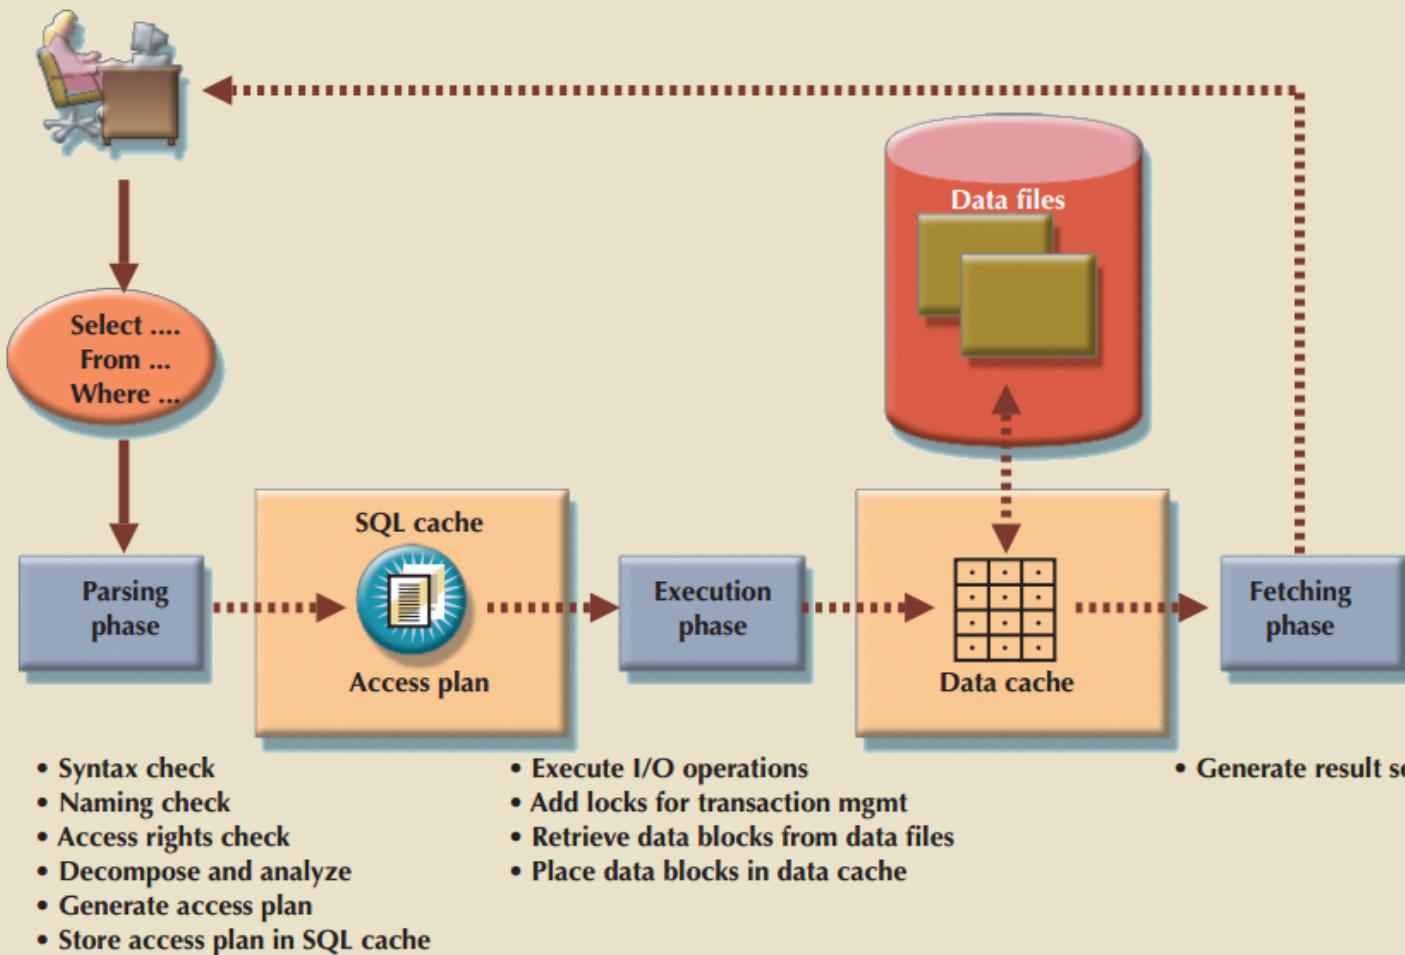
- DBMS executes the SQL query using the chosen execution plan

## Fetching

- DBMS fetches the data and sends the result set back to the client

# Figure 11.2 – Query Processing

FIGURE 11.2 QUERY PROCESSING



# SQL Parsing Phase

- Query is broken down into smaller units
- Original SQL query transformed into slightly different version of original SQL code which is fully equivalent and more efficient
- **Query optimizer:** Analyzes SQL query and finds most efficient way to access data
- **Access plans:** DBMS-specific and translate client's SQL query into a series of complex I/O operations
- If access plan already exists for query in SQL cache, DBMS reuses it
  - If not, optimizer evaluates various plans and chooses one to be placed in SQL cache for use

# SQL Execution Phase

- All I/O operations indicated in the access plan are executed
  - Locks are acquired
  - Data are retrieved and placed in data cache
  - Transaction management commands are processed

# SQL Fetching Phase

- Rows of resulting query result set are returned to client
  - DBMS may use temporary table space to store temporary data
  - Database server coordinates the movement of the result set rows from the server cache to the client cache

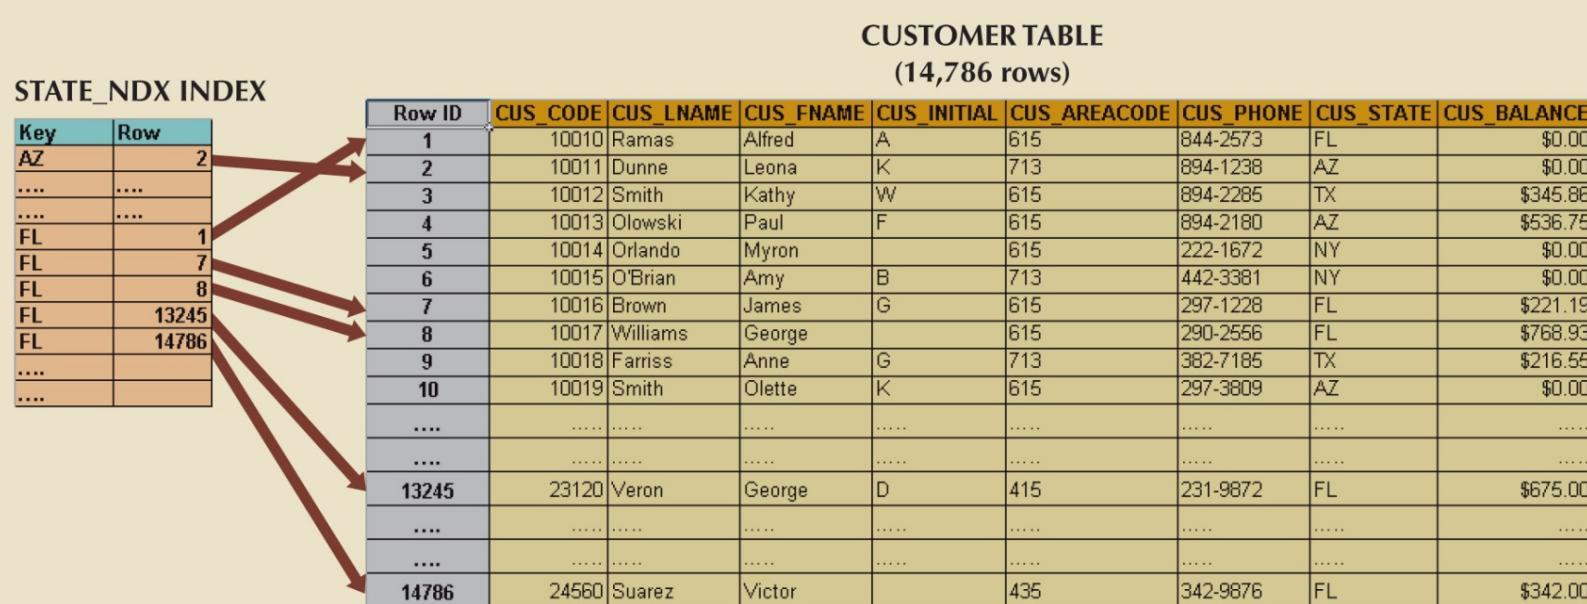
# Query Processing Bottlenecks

- Delay introduced in the processing of an I/O operation that slows the system
- Caused by the:
  - CPU
  - RAM
  - Hard disk
  - Network
  - Application code

# Indexes and Query Optimization

- Indexes
  - Help speed up data access
  - Facilitate searching, sorting, using aggregate functions, and join operations
  - Ordered set of values that contain the index key and pointers
  - More efficient than a full table scan

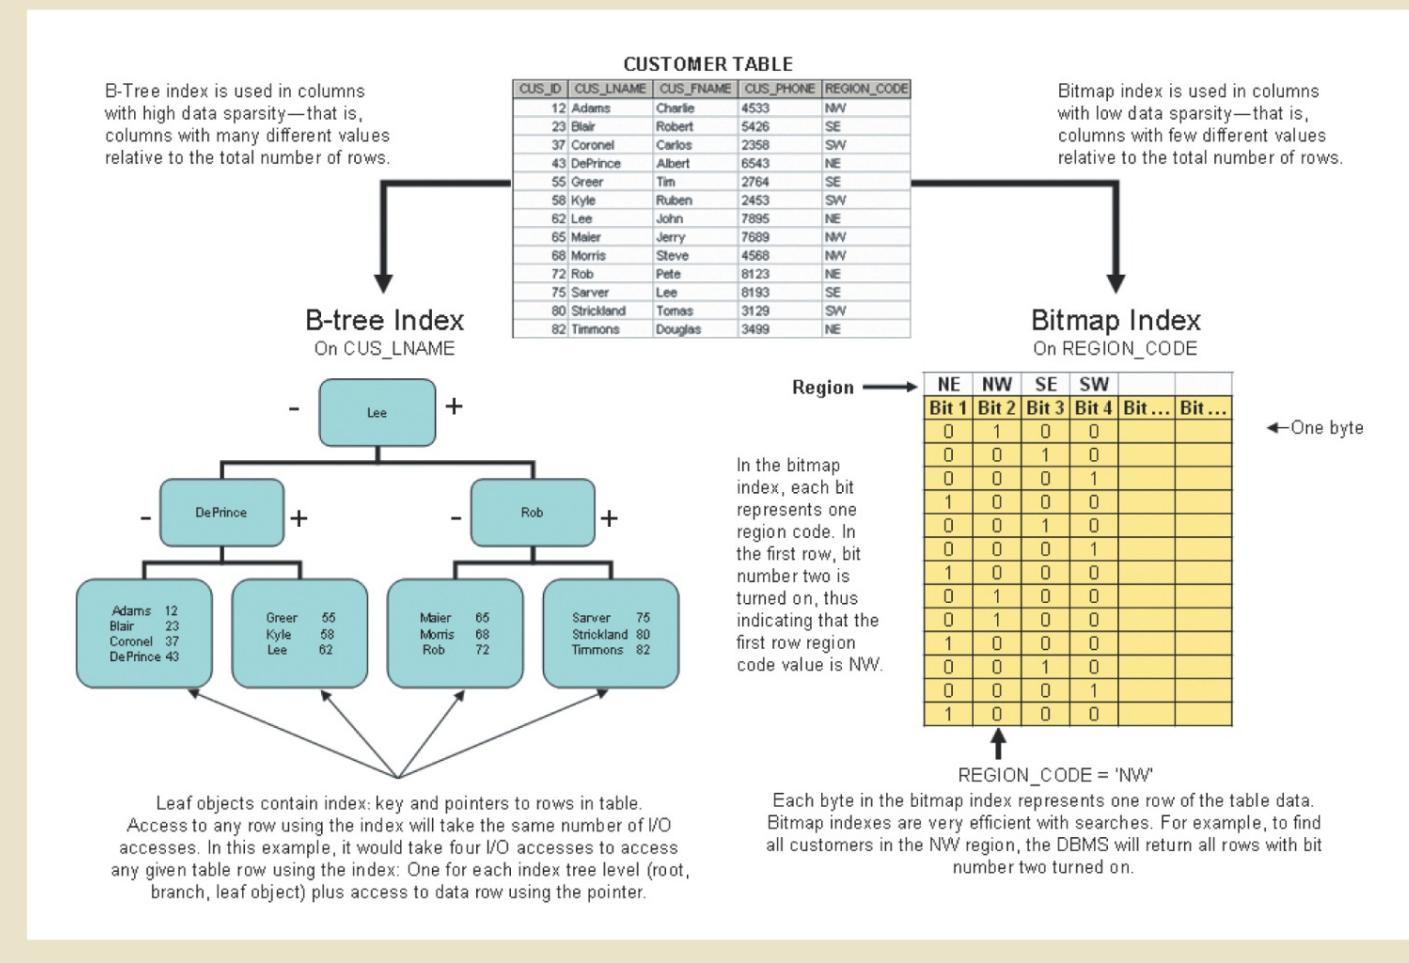
**FIGURE 11.3 INDEX REPRESENTATION FOR THE CUSTOMER TABLE**



# Indexes and Query Optimization

- **Data sparsity:** Number of different values a column could have
- Data structures used to implement indexes:
  - **Hash indexes**
  - **B-tree indexes**
  - **Bitmap indexes**
- DBMSs determine best type of index to use

FIGURE 11.4 B-TREE AND BITMAP INDEX REPRESENTATION



# Optimizer Choices

- **Rule-based optimizer:** Uses preset rules and points to determine the best approach to execute a query
- **Cost-based optimizer:** Uses algorithms based on statistics about objects being accessed to determine the best approach to execute a query

TABLE 11.4

## COMPARING ACCESS PLANS AND I/O COSTS

PLAN	STEP	OPERATION	I/O OPERATIONS	I/O COST	RESULTING SET ROWS	TOTAL I/O COST
<b>A</b>	A1	Cartesian product (PRODUCT, VENDOR)	7,000 + 300	7,300	2,100,000	7,300
	A2	Select rows in A1 with matching vendor codes	2,100,000	2,100,000	7,000	2,107,300
	A3	Select rows in A2 with V_STATE = 'FL'	7,000	7,000	1,000	<b>2,114,300</b>
<b>B</b>	B1	Select rows in VENDOR with V_STATE = 'FL'	300	300	10	300
	B2	Cartesian Product (PRODUCT, B1)	7,000 + 10	7,010	70,000	7,310
	B3	Select rows in B2 with matching vendor codes	70,000	70,000	1,000	<b>77,310</b>

# Using Hints to Affect Optimizer Choices

- Optimizer might not choose the best execution plan
  - Makes decisions based on existing statistics, which might be old
  - Might choose less-efficient decisions
- **Optimizer hints:** Special instructions for the optimizer, embedded in the SQL command text

# Table 11.5 - Optimizer Hints

TABLE 11.5

## OPTIMIZER HINTS

HINT	USAGE
ALL_ROWS	Instructs the optimizer to minimize the overall execution time—that is, to minimize the time needed to return all rows in the query result set. This hint is generally used for batch mode processes. For example: <pre>SELECT /*+ ALL_ROWS */ *       FROM PRODUCT      WHERE P_QOH &lt; 10;</pre>
FIRST_ROWS	Instructs the optimizer to minimize the time needed to process the first set of rows—that is, to minimize the time needed to return only the first set of rows in the query result set. This hint is generally used for interactive mode processes. For example: <pre>SELECT /*+ FIRST_ROWS */ *       FROM PRODUCT      WHERE P_QOH &lt; 10;</pre>
INDEX(name)	Forces the optimizer to use the P_QOH_NDX index to process this query. For example: <pre>SELECT /*+ INDEX(P_QOH_NDX) */ *       FROM PRODUCT      WHERE P_QOH &lt; 10</pre>

# SQL Performance Tuning

- Evaluated from client perspective
  - Most current relational DBMSs perform automatic query optimization at the server end
  - Most SQL performance optimization techniques are DBMS-specific and thus rarely portable
- Majority of performance problems are related to poorly written SQL code

# Index Selectivity

- Measure of the likelihood that an index will be used in query processing
- Indexes are used when a subset of rows from a large table is to be selected based on a given condition
- Index cannot always be used to improve performance
- **Function-based index:** Based on a specific SQL function or expression

# Conditional Expressions

- Expressed within WHERE or HAVING clauses of a SQL statement
  - Restricts the output of a query to only rows matching conditional criteria
- Guidelines to write efficient conditional expressions in SQL code
  - Use simple columns or literals as operands
  - Numeric field comparisons are faster than character, date, and NULL comparisons

# Conditional Expressions

- Equality comparisons are faster than inequality comparisons
- Transform conditional expressions to use literals
- Write equality conditions first when using multiple conditional expressions
- When using multiple AND conditions, write the condition most likely to be false first
- When using multiple OR conditions, put the condition most likely to be true first
- Avoid the use of NOT logical operator

# Table 11.6 – Conditional Criteria

TABLE 11.6

## CONDITIONAL CRITERIA

OPERAND1	CONDITIONAL OPERATOR	OPERAND2
P_PRICE	>	10.00
V_STATE	=	FL
V_CONTACT	LIKE	Smith%
P_QOH	>	P_MIN * 1.10

# Query Formulation

- Identify what columns and computations are required
- Identify source tables
- Determine how to join tables
- Determine what selection criteria are needed
- Determine the order in which to display the output

# DBMS Performance Tuning

- Managing DBMS processes in primary memory and the structures in physical storage
- DBMS performance tuning at server end focuses on setting parameters used for:
  - Data cache
  - SQL cache
  - Sort cache
  - Optimizer mode
- **In-memory database:** Store large portions of the database in primary storage

# DBMS Performance Tuning

- Recommendations for physical storage of databases:
  - Use **RAID** (Redundant Array of Independent Disks) to provide a balance between performance improvement and fault tolerance
  - Minimize disk contention
  - Put high-usage tables in their own table spaces
  - Assign separate data files in separate storage volumes for indexes, system, and high-usage tables

# DBMS Performance Tuning

- Take advantage of the various table storage organizations in the database
  - **Index-organized table or clustered index table:** Stores the end-user data and the index data in consecutive locations in permanent storage
- Partition tables based on usage
- Use denormalized tables where appropriate
- Store computed and aggregate attributes in tables

# Query Optimization Example

FIGURE 11.5 INITIAL EXPLAIN PLAN

The screenshot shows a SQL Plus window with the following content:

```
SQL> ANALYZE TABLE QOUENDOR COMPUTE STATISTICS;
Table analyzed.

SQL> EXPLAIN PLAN FOR SELECT * FROM QOUENDOR WHERE U_NAME LIKE 'B%' ORDER BY U_AREACODE;
Explained.

SQL> SELECT * FROM TABLE(DBMS_XPLAN.DISPLAY);
PLAN_TABLE_OUTPUT
```

---

Plan hash value: 1800591659

Id	Operation	Name	Rows	Bytes	Cost	CPU	Time
0	SELECT STATEMENT		1	38	3	<0>	00:00:01
1	SORT ORDER BY		1	38	3	<0>	00:00:01
2	TABLE ACCESS FULL	QOUENDOR	1	38	3	<0>	00:00:01

---

Predicate Information (identified by operation id):

```
PLAN_TABLE_OUTPUT
```

---

```
2 - filter<"U_NAME" LIKE 'B%'>
14 rows selected.
```

```
SQL> _
```

FIGURE 11.6 EXPLAIN PLAN AFTER INDEX ON V\_AREACODE

The screenshot shows a SQL Plus window with the following content:

```
SQL> CREATE INDEX QOU_NDX1 ON QOUENDOR(V_AREACODE);
Index created.

SQL> ANALYZE TABLE QOUENDOR COMPUTE STATISTICS;
Table analyzed.

SQL> EXPLAIN PLAN FOR SELECT * FROM QOUENDOR WHERE V_NAME LIKE 'B%' ORDER BY V_AREACODE;
Explained.

SQL> SELECT * FROM TABLE(DBMS_XPLAN.DISPLAY);
PLAN_TABLE_OUTPUT
```

---

Plan hash value: 641227332

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time
0	SELECT STATEMENT		1	38	2 (0)	00:00:01
1	TABLE ACCESS BY INDEX ROWID	QOUENDOR	1	38	2 (0)	00:00:01
2	INDEX FULL SCAN	QOU_NDX1	15		1 (0)	00:00:01

---

Predicate Information (identified by operation id):

```
PLAN_TABLE_OUTPUT
```

---

-----  
1 - filter("V\_NAME" LIKE 'B%')  
14 rows selected.

SQL> \_

FIGURE 11.7 EXPLAIN PLAN AFTER INDEX ON V\_NAME

```
SQL> CREATE INDEX QOU_NDX2 ON QOUENDAR(U_NAME);
Index created.

SQL> ANALYZE TABLE QOUENDAR COMPUTE STATISTICS;
Table analyzed.

SQL> EXPLAIN PLAN FOR SELECT * FROM QOUENDAR WHERE U_NAME LIKE 'B%' ORDER BY U_AREACODE;
Explained.

SQL> SELECT * FROM TABLE(DBMS_XPLAN.DISPLAY);
PLAN_TABLE_OUTPUT

Plan hash value: 641227332

| Id  | Operation          | Name    | Rows  | Bytes | Cost (%CPU)| Time     |
|---|---|---|---|---|---|---|
| 0  | SELECT STATEMENT   |         |       |       | 2  (0) | 00:00:01 |
| * 1 | TABLE ACCESS BY INDEX ROWID | QOUENDAR | 1    | 38   | 2  (0) | 00:00:01 |
| 2  |   INDEX FULL SCAN  | QOU_NDX1 | 15   | 38   | 1  (0) | 00:00:01 |

Predicate Information (identified by operation id):
PLAN_TABLE_OUTPUT

1 - filter("U_NAME" LIKE 'B%')
14 rows selected.

SQL> _
```

FIGURE 11.8 ACCESS PLAN USING INDEX ON V\_NAME

The screenshot shows a SQL Plus window with the following content:

```
SQL> EXPLAIN PLAN FOR SELECT U_NAME, P_CODE FROM QO_VENDOR U, QO_PRODUCT P
  2 WHERE U.U_CODE = P.U_CODE AND U_NAME = 'ORDUA, Inc.';
Explained.
SQL> SELECT * FROM TABLE(DBMS_XPLAN.DISPLAY);
PLAN_TABLE_OUTPUT
```

**Plan hash value: 4146133223**

#	Id	Operation	Name	Rows	Bytes	Cost	%CPU	Time
	0	SELECT STATEMENT		2	62	5	<0>	00:00:01
*	1	HASH JOIN		2	62	5	<0>	00:00:01
	2	TABLE ACCESS BY INDEX ROWID BATCHED	QO_VENDOR	1	17	2	<0>	00:00:01
*	3	INDEX RANGE SCAN	QO_NDX2	1	1	1	<0>	00:00:01
*	4	TABLE ACCESS FULL	QO_PRODUCT	14	196	3	<0>	00:00:01

**PLAN\_TABLE\_OUTPUT**

**Predicate Information (identified by operation id):**

```
1 - access("U"."U_CODE"="P"."U_CODE")
3 - access("U_NAME"='ORDUA, Inc.')
4 - filter("P"."U_CODE" IS NOT NULL)
```

18 rows selected.

SQL> \_

FIGURE 11.9 ACCESS PLAN USING FUNCTIONS ON INDEXED COLUMNS

```
SQL> EXPLAIN PLAN FOR SELECT U_NAME, P_CODE FROM QO_VENDOR U, QO_PRODUCT P
  2          WHERE U.U_CODE = P.U_CODE AND UPPER(U_NAME) = 'ORDUA, INC.';
Explained.
SQL> SELECT * FROM TABLE(DBMS_XPLAN.DISPLAY);
PLAN_TABLE_OUTPUT
Plan hash value: 1347990970

| Id | Operation           | Name      | Rows  | Bytes | Cost (%CPU)| Time     |
|---|---|---|---|---|---|---|
| 0 | SELECT STATEMENT   |           | 1    | 31   | 5  (0) | 00:00:01 |
|* 1 | HASH JOIN           |           | 1    | 31   | 5  (0) | 00:00:01 |
|* 2 |   VIEW              | index$_join$_001 | 1    | 17   | 2  (0) | 00:00:01 |
|* 3 |   HASH JOIN          |           | 1    | 17   | 1  (0) | 00:00:01 |
|* 4 |   INDEX FAST FULL SCAN | QOU_NDX2 | 1    | 17   | 1  (0) | 00:00:01 |
| 5 |   INDEX FAST FULL SCAN | SYS_C0058569 | 1    | 17   | 1  (0) | 00:00:01 |
PLAN_TABLE_OUTPUT
|* 6 | TABLE ACCESS FULL   | QOPRODUCT | 14   | 196  | 3  (0) | 00:00:01 |

Predicate Information (identified by operation id):
1 - access("U"."U_CODE"="P"."U_CODE")
3 - access(ROWID=ROWID)
4 - filter(UPPER("U_NAME")='ORDUA, INC.')
6 - filter("P"."U_CODE" IS NOT NULL)

21 rows selected.

SQL> _
```

FIGURE 11.10 FIRST EXPLAIN PLAN: AGGREGATE FUNCTION ON A NON-INDEXED COLUMN

The screenshot shows a SQL Plus window with the following content:

```
SQL> ANALYZE TABLE QOPRODUCT COMPUTE STATISTICS;
Table analyzed.

SQL> EXPLAIN PLAN FOR SELECT MAX(P_PRICE) FROM QOPRODUCT;
Explained.

SQL> SELECT * FROM TABLE(DBMS_XPLAN.DISPLAY);
PLAN_TABLE_OUTPUT
-----
Plan hash value: 1624837700

| Id  | Operation          | Name      | Rows   | Bytes  | Cost  | %CPU | Time      |
| 0   | SELECT STATEMENT   |           |        |        |       |       |            |
| 1   |   SORT AGGREGATE   |           | 1     | 4      | 3     | 0    | <0> :00:01 |
| 2   |     TABLE ACCESS FULL| QOPRODUCT | 16    | 64     | 3     | 0    | <0> :00:01 |

9 rows selected.

SQL> _
```

The window has a blue header bar with the title "SQL Plus". The main area contains the SQL commands and their output. The output shows an explain plan with three operations: a SELECT STATEMENT, a SORT AGGREGATE, and a TABLE ACCESS FULL. The TABLE ACCESS FULL operation is for the QOPRODUCT table, which has 16 rows and 64 bytes. The SORT AGGREGATE operation has one row and 4 bytes. The total cost is 3, and the execution time is 00:00:01.

FIGURE 11.11 SECOND EXPLAIN PLAN: AGGREGATE FUNCTION ON AN INDEXED COLUMN

The screenshot shows a SQL Plus session window with the following content:

```
SQL> CREATE INDEX QOP_NDX2 ON QOPRODUCT(P_PRICE);
Index created.

SQL> ANALYZE TABLE QOPRODUCT COMPUTE STATISTICS;
Table analyzed.

SQL> EXPLAIN PLAN FOR SELECT MAX(P_PRICE) FROM QOPRODUCT;
Explained.

SQL> SELECT * FROM TABLE(DBMS_XPLAN.DISPLAY);
PLAN_TABLE_OUTPUT
```

---

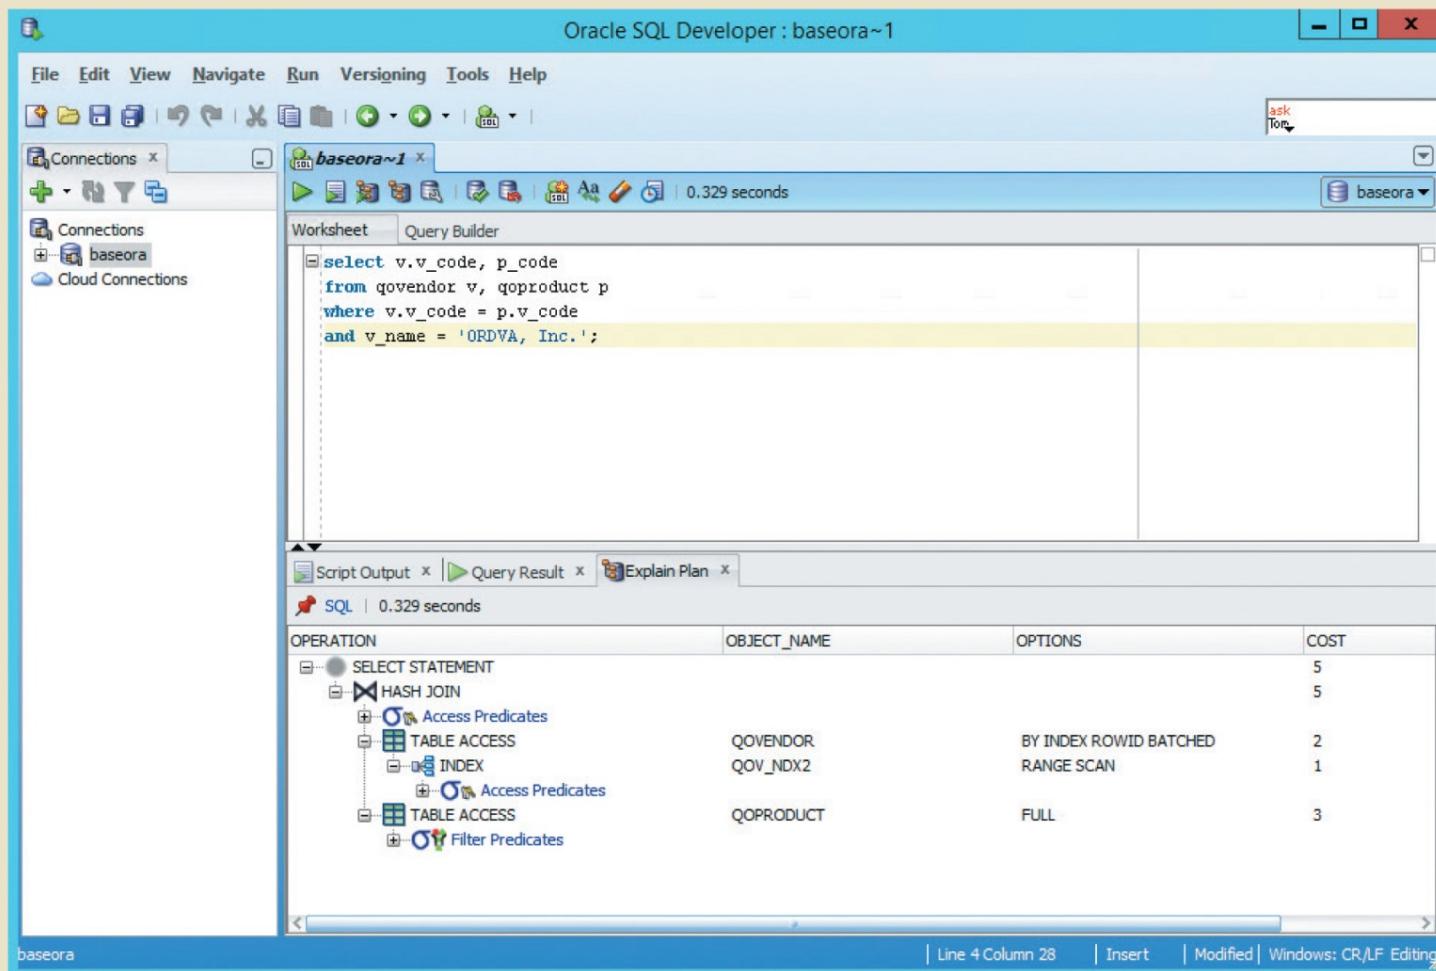
Plan hash value: 3880272194

Id	Operation	Name	Rows	Bytes	Cost	%CPU	Time
0	SELECT STATEMENT		1	4	1	<0>	00:00:01
1	SORT AGGREGATE		1	4			
2	INDEX FULL SCAN (MIN/MAX)	QOP_NDX2	1	4	1	<0>	00:00:01

9 rows selected.

SQL> \_

FIGURE 11.12 ORACLE TOOLS FOR QUERY OPTIMIZATION



**FIGURE 11.13 MYSQL TOOLS FOR QUERY OPTIMIZATION**

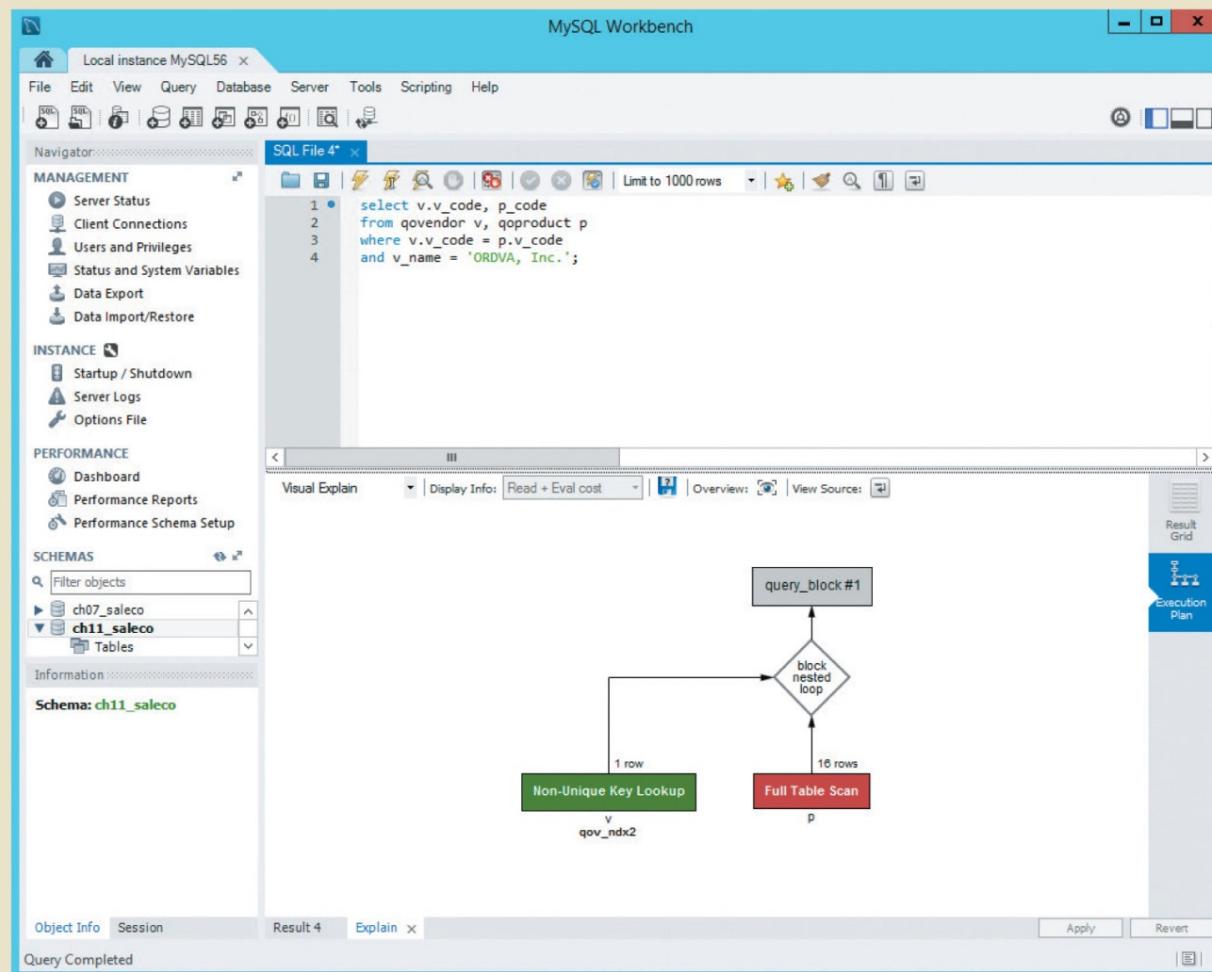


FIGURE 11.14 MICROSOFT SQL SERVER TOOLS FOR QUERY OPTIMIZATION

